

Medical OCR API — Assignment Summary Report

Project: Medical OCR API (Home Assessment – OCR Endpoint Challenge)

Author: Shou Yichen

Date: 2026-08-05 **Tech Stack:** FastAPI + Tesseract (pytesseract) + OpenAI

LLM + OpenCV **Test Environment:** macOS 26.6 (arm64) on Apple M1 Pro (8

cores); Python 3.9.6 (project .venv); pytest 8.4.2; Tesseract 5.5.3


(/opt/homebrew/bin/tesseract); automated tests: 5 passed, 5 warnings

Table of Contents

1. [Pipeline Architecture](#)
 2. [API Input & Output Results](#)
 3. [Error Handling Verification](#)
 4. [Performance Summary](#)
 5. [Key Design Decisions](#)
-

1. Pipeline Architecture

1.1 Architecture Diagram



1.2 Web UI

The project includes a browser-based drag-and-drop interface for easy testing:

Medical OCR API
Upload a medical document — referral letter, medical certificate, or receipt — and get structured JSON.

Drag & drop a file here, or click to browse
receipt.pdf
Supported: PDF, JPEG, PNG, TIFF (max 20 MB)

Upload & Analyze

✓ Processing completed.

Result Copy JSON

Document: receipt Total: 25.5435s OCR: 4.8685s Classify: 20.5311s Extract: 20.5311s

```
{
  "claimant_name": "JOHN DOE",
  "claimant_address": "123 SAMPLE ST #01-01",
  "claimant_date_of_birth": null,
  "provider_name": "Raffles Medical",
  "tax_amount": 3,
  "total_amount": 49
}
```

1.3 Online Demo

The deployed online demo is available at:

- <https://medical-ocr.greensky-c0f2aeba.eastasia.azurecontainerapps.io>

1.4 Online Test Cases

The following curl requests use the repository sample files with relative paths:

```
# Referral Letter
curl --fail-with-body -X POST "https://medical-ocr.greensky-
c0f2aeba.eastasia.azurecontainerapps.io/ocr" \
-F "file=@Example/referral_letter.pdf"
```

```
# Medical Certificate
curl --fail-with-body -X POST "https://medical-ocr.greensky-
c0f2aeba.eastasia.azurecontainerapps.io/ocr" \
-F "file=@Example/medical_certificate.pdf"

# Receipt
curl --fail-with-body -X POST "https://medical-ocr.greensky-
c0f2aeba.eastasia.azurecontainerapps.io/ocr" \
-F "file=@Example/receipt.pdf"
```

2. API Input & Output Results

2.1 Referral Letter

Input: Example/referral_letter.pdf

Healthway Screening @ Centrepont

176 Orchard Road, #06-03/04
The Centrepont, Singapore 238843



JOHN DOE

ID:

ID: S1234567A IC: S9876543Z

12/09/2022

01/01/2025

Dear dr.

Thank you for seeing the above pt. He has multiple moles and freckles on his body and reports that there has been a change in size and shape for some of them. Kindly assist to perform a mole check.

Kind regards,

DR. SAMPLE LEE

Healthway Screening @ Centrepont
176 Orchard Road,
#06-03/04, The Centrepont
Singapore 238843
Tel: 6811 8686 Fax: 6733 5951

Request:

```
POST /ocr
Content-Type: multipart/form-data
file=@Example/referral_letter.pdf
```

Output (200 OK):

```
{
  "message": "Processing completed.",
  "result": {
    "document_type": "referral_letter",
    "total_time": 22.7419,
    "ocr_time": 6.8492,
    "classification_time": 15.6972,
    "extraction_time": 15.6972,
    "finalJson": {
      "claimant_name": "JOHN DOE",
      "provider_name": "Healthway Screening @ Centrepont",
      "signature_presence": false,
      "total_amount_paid": null,
      "total_approved_amount": null,
      "total_requested_amount": null
    }
  }
}
```

Field verification:

Field	Expected	Actual	✓/X
claimant_name	Patient Name	"JOHN DOE"	✓
provider_name	No "Fullerton Health"	"Healthway Screening @ Centrepont"	✓
signature_presence	true/false	false	✓
total_amount_paid	Integer	null (not present in doc)	✓

total_approved_amount	Integer	null (not present in doc)	✓
total_requested_amount	Integer	null (not present in doc)	✓

Note on @ symbol: The LLM prompt includes OCR correction for @ misread as 0, o, a, or at. The provider name "Healthway Screening @ Centrepont" was correctly extracted.

2.2 Medical Certificate

Input: Example/medical_certificate.pdf

12/1/22, 11:10 AM

Digital Medical Certificate



REG NO.: REG2025010101

NAME: JOHN DOE

NRIC / FIN / PASSPORT: S1234567A

This is to certify that the above-named is unfit for duty for a period of 1 days from 30-Nov-2022 to 30-Nov-2022.

TYPE OF MEDICAL CERTIFICATE:

- ☐ Hospitalisation Leave
- ☒ Outpatient Sick Leave
- ☐ Maternity Leave
- ☐ Sterilisation Leave
- ☐ Time Chit

Admitted on: _____ Discharged on: _____

Delivered on: _____

Operated on: _____

Time In: _____ Time Out: _____

This certificate is not valid for absence from court attendance.

Fit for light duty from N.A. to N.A..

COMMENTS:

HOSPITAL/CLINIC	WARD	NAME/DESIGNATION/MCR NO
Minmed Health Screeners	NA	
	DATE	
	30-Nov-2022	

This medical certificate is electronically generated. No signature is required.

Request:

```
POST /ocr
Content-Type: multipart/form-data
file=@Example/medical_certificate.pdf
```

Output (200 OK):

```
{
  "message": "Processing completed.",
  "result": {
    "document_type": "medical_certificate",
    "total_time": 21.4516,
    "ocr_time": 2.7371,
    "classification_time": 18.5032,
    "extraction_time": 18.5032,
    "finalJson": {
      "claimant_name": "JOHN DOE",
      "claimant_address": null,
      "claimant_date_of_birth": null,
      "diagnosis_name": null,
      "discharge_date_time": null,
      "icd_code": null,
      "provider_name": "Minmed Health Screeners",
      "submission_date_time": null,
      "date_of_mc": "30/11/2022",
      "mc_days": 1
    }
  }
}
```

Field verification:

Field	Expected	Actual	✓/✗
claimant_name	Claimant Name	"JOHN DOE"	✓
claimant_address	Address	null (not in doc)	✓
claimant_date_of_birth	DD/MM/YYYY	null (not in doc)	✓
diagnosis_name	Diagnosis	null (not in doc)	✓

discharge_date_time	DD/MM/YYYY	null (not in doc)	✓
icd_code	ICD code	null (not in doc)	✓
provider_name	No "Fullerton Health"	"Minmed Health Screeners"	✓
submission_date_time	DD/MM/YYYY	null (not in doc)	✓
date_of_mc	DD/MM/YYYY	"30/11/2022"	✓
mc_days	Integer	1	✓

2.3 Receipt

Input: Example/receipt.pdf

Raffles Medical
Your Trusted Partner for Health

TAX INVOICE

GST123456Z

VIS: 20230104-000
VISIT DATE/TIME: 04-JAN-2023 12:25PM
BILL DATE: 04-JAN-2023

INV: 20230104-000
PATIENT ID NO.: [REDACTED]

PAY BY: SELF
JOHN DOE

123 SAMPLE ST. #01-01

DESCRIPTION	QTY	SE	SS
CONSULTATION			28.00
PHARMACEUTICAL			
DEXTRAMETHORPHAN 15MG/5ML SYR	80.0	5.47	
TROSPRONIT MAX LOZENGES	8.0	3.16	
PRACTICE COST			8.63
PRACTICE COST	1.0	11.00	
SUB-TOTAL			11.00
TOTAL CHARGES BEFORE GST			45.63
GST @ 8%			3.65
TOTAL CHARGES AFTER GST			49.28
LESS ROUNDING ADJUSTMENT			(0.03)
TOTAL AMOUNT PAID			(49.25)
VISA ****1234			
TOTAL BALANCE DUE			0.00

Raffles Medical
21 Choa Chu Kang North 6
New Rise Point #01-02
Singapore 680576
Tel: (65) 6534 3132 Fax: (65) 6534 3510

21 CHOA CHU KANG NORTH 6 #01-02 NEW RISE POINT SINGAPORE 680576 T: 65343132
Raffles Medical Group Ltd | Company Registration No: 198901967K | GST Registration No: M9-0200467-N

Request:

```
POST /ocr
Content-Type: multipart/form-data
file=@Example/receipt.pdf
```

Output (200 OK):

```
{
  "message": "Processing completed.",
  "result": {
    "document_type": "receipt",
    "total_time": 26.1796,
    "ocr_time": 5.476,
    "classification_time": 20.5064,
    "extraction_time": 20.5064,
    "finalJson": {
      "claimant_name": "JOHN DOE",
      "claimant_address": "123 SAMPLE ST #01-01",
      "claimant_date_of_birth": null,
      "provider_name": "Raffles Medical",
      "tax_amount": 3,
      "total_amount": 49
    }
  }
}
```

Field verification:

Field	Expected	Actual	✓/X
claimant_name	Claimant Name	"JOHN DOE"	✓
claimant_address	Address	"123 SAMPLE ST #01-01"	✓
claimant_date_of_birth	DD/MM/YYYY	null (not in doc)	✓
provider_name	No "Fullerton Health"	"Raffles Medical"	✓

tax_amount	Integer	3	✓
total_amount	Integer	49	✓

Note on amounts: tax_amount: 3 and total_amount: 49 are correctly stripped of currency symbols, decimal separators, and returned as integers. The #01-01 address format was correctly OCR-corrected.

3. Error Handling Verification

Test Case	Expected	Actual	✓/✗
No file attached	400 file_missing	400	✓
Invalid MIME type	400 file_missing	400	✓
Empty filename	422 (FastAPI layer)	422	✓
Unsupported document	422 unsupported_document_type	422	✓
Upstream LLM failure	500 internal_server_error	500	✓

Automated tests: 5/5 passing (run in project virtualenv).

```
$ pytest tests/test_api.py -v
tests/test_api.py::TestAPI::test_health_check PASSED
tests/test_api.py::TestAPI::test_missing_file PASSED
tests/test_api.py::TestAPI::test_invalid_mime PASSED
tests/test_api.py::TestAPI::test_unsupported_document PASSED
```

```
tests/test_api.py::TestAPI::test_empty_filename PASSED
===== 5 passed, 5 warnings in 0.09s
=====
```

4. Performance Summary

Document	OCR Time	LLM Time	Total Time	Status
referral_letter.pdf	6.85s	15.70s	22.74s	200
medical_certificate.pdf	2.74s	18.50s	21.45s	200
receipt.pdf	5.48s	20.51s	26.18s	200

Breakdown: OCR (Tesseract/pytesseract) takes ~3–7s. LLM inference takes ~15–21s (network latency to remote API). Total end-to-end: ~21–26s per document.

5. Key Design Decisions

5.1 Single LLM Call for Classification + Extraction

Instead of a two-step pipeline (classify → extract), the LLM performs both in a single API call. The system prompt includes full JSON templates for all three document types, and the LLM copies the correct template, fills values, and returns valid JSON. This reduces latency and cost by 50%.

5.2 Multi-Feature Signature Detection

Signature detection uses a 4-dimensional weighted scoring system:

- **Area dispersion** (CV of contour areas)
- **Stroke thinness** (aspect ratio of bounding rects)
- **Curvature** (circularity + solidity)
- **Ink density** (fill ratio 2%–30%)

Stamps/seals and barcodes are explicitly excluded before scoring via Hough Circle Transform and Canny edge projection analysis.

5.3 OCR Error Correction (Two-Layer Defense)

Layer	Mechanism
Prompt	LLM is instructed to fix: @ misread, name trailing debris, Singapore unit numbers
Post-processing	<code>_coerce_value()</code> regex fallback for <code>claimant_name</code> and <code>claimant_address</code>

5.4 Extensibility

Adding a new document type requires only:

6. Add to `DocumentType` enum
7. Add JSON template to LLM system prompt

8. Optionally add type-specific post-processing (like signature detection)

No new extractor classes, regex rules, or classifier training needed.

Appendix

A. Quick Start

```
# Install dependencies
pip install -r requirements.txt
sudo apt-get install -y poppler-utils

# Install Tesseract OCR binary (required for pytesseract)
# macOS (Homebrew):
#   brew install tesseract
# Ubuntu / Debian:
#   sudo apt-get install -y tesseract-ocr tesseract-ocr-eng
# Windows (Chocolatey):
#   choco install tesseract

# Configure LLM
cp .env.example .env
# Edit .env: LLM_API_KEY, LLM_MODEL, LLM_BASE_URL

# Run
uvicorn app.main:app --host 0.0.0.0 --port 8000

# Test
curl -X POST http://localhost:8000/ocr -F "file=@Example/referral_letter.pdf"
```

B. Project Structure

```
medical-ocr-api/
├── app/
│   ├── main.py           # FastAPI entry point
│   ├── config.py         # Environment config
│   ├── api/ocr.py        # POST /ocr endpoint
│   ├── models/response.py # Pydantic models
│   ├── pipeline/
│   │   ├── preprocessor.py # PDF → Image
│   │   ├── ocr_engine.py   # Tesseract (pytesseract)
│   │   ├── normalizer.py   # Text cleanup
│   │   └── ocr_types.py    # Dataclasses
│   ├── extraction/
│   │   ├── llm_extractor.py # LLM classify + extract
│   │   └── signature.py     # Handwriting detection
│   └── static/index.html   # Web UI
```

```
└─ utils/
└─ Example/          # Sample documents
└─ tests/            # pytest
└─ doc/              # Documentation
└─ Dockerfile
└─ docker-compose.yml
└─ README.md
```